

```
In [1]: #Name : Akanksha Chandranohan Giri
#Roll no : 41
#Section : 3A
#Date : 05/10/2024
```

```
In [1]: #Aim : To perform operation on KNN (K Nearest Neighbor)
```

Importing Libraries

```
In [7]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings('ignore')
```

```
In [8]: import os
```

```
In [9]: os.getcwd()
```

```
Out[9]: 'C:\\Users\\HP'
```

```
In [10]: os.chdir("C:\\Users\\HP\\Desktop")
```

```
In [11]: df=pd.read_csv("framingham.csv")
```

```
In [12]: #The "Framingham" heart disease dataset includes over 4,240 records, 15 attributes
#The goal of the dataset is to predict whether the patient has 10-year risk of
```

```
In [13]: df.head()
```

```
Out[13]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diab
0	1	39	4.0	0	0.0	0.0	0	0	
1	0	46	2.0	0	0.0	0.0	0	0	
2	1	48	1.0	1	20.0	0.0	0	0	
3	0	61	3.0	1	30.0	0.0	0	1	
4	0	46	3.0	1	23.0	0.0	0	0	

```
In [14]: df.describe()
```

```
Out[14]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentS
count	4238.000000	4238.000000	4133.000000	4238.000000	4209.000000	4185.000000	4238.0
mean	0.429212	49.584946	1.978950	0.494101	9.003089	0.029630	0.0
std	0.495022	8.572160	1.019791	0.500024	11.920094	0.169584	0.0
min	0.000000	32.000000	1.000000	0.000000	0.000000	0.000000	0.0
25%	0.000000	42.000000	1.000000	0.000000	0.000000	0.000000	0.0
50%	0.000000	49.000000	2.000000	0.000000	0.000000	0.000000	0.0

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentS
75%	1.000000	56.000000	3.000000	1.000000	20.000000	0.000000	0.0
max	1.000000	70.000000	4.000000	1.000000	70.000000	1.000000	1.0

In [15]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4238 entries, 0 to 4237
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  4238 non-null   int64
1   age                   4238 non-null   int64
2   education             4133 non-null   float64
3   currentSmoker         4238 non-null   int64
4   cigsPerDay            4209 non-null   float64
5   BPMeds                4185 non-null   float64
6   prevalentStroke       4238 non-null   int64
7   prevalentHyp          4238 non-null   int64
8   diabetes              4238 non-null   int64
9   totChol               4188 non-null   float64
10  sysBP                 4238 non-null   float64
11  diaBP                 4238 non-null   float64
12  BMI                   4219 non-null   float64
13  heartRate             4237 non-null   float64
14  glucose               3850 non-null   float64
15  TenYearCHD            4238 non-null   int64
dtypes: float64(9), int64(7)
memory usage: 529.9 KB
```

In [16]: `df.isna().sum()`

```
Out[16]: male                0
age                  0
education            105
currentSmoker        0
cigsPerDay           29
BPMeds               53
prevalentStroke      0
prevalentHyp         0
diabetes              0
totChol              50
sysBP                0
diaBP                0
BMI                  19
heartRate            1
glucose              388
TenYearCHD           0
dtype: int64
```

In []: *#Since, only a few rows have null values in them, we are only removing those*
#df = df.dropna(subset=['heartRate', 'BMI', 'cigsPerDay', 'totChol', 'BPMeds'])

In [17]: `df`

```
Out[17]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	d
0	1	39	4.0	0	0.0	0.0	0	0	

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	d
1	0	46	2.0	0	0.0	0.0	0	0	
2	1	48	1.0	1	20.0	0.0	0	0	
3	0	61	3.0	1	30.0	0.0	0	1	
4	0	46	3.0	1	23.0	0.0	0	0	
...	
4233	1	50	1.0	1	1.0	0.0	0	1	
4234	1	51	3.0	1	43.0	0.0	0	0	
4235	0	48	2.0	1	20.0	NaN	0	0	
4236	0	44	1.0	1	15.0	0.0	0	0	
4237	0	52	2.0	0	0.0	0.0	0	0	

4238 rows × 16 columns

Missing Value Treatment

Since, 'glucose' and 'education' columns had a significant amount of null values, so we replaced them with the mean of values for their respective columns

```
In [18]: df['glucose'].fillna(value = df['glucose'].mean(),inplace=True)

In [19]: df['education'].fillna(value = df['education'].mean(),inplace=True)

In [20]: df['heartRate'].fillna(value = df['heartRate'].mean(),inplace=True)

In [21]: df['BMI'].fillna(value = df['BMI'].mean(),inplace=True)

In [22]: df['cigsPerDay'].fillna(value = df['cigsPerDay'].mean(),inplace=True)

In [23]: df['totChol'].fillna(value = df['totChol'].mean(),inplace=True)

In [24]: df['BPMeds'].fillna(value = df['BPMeds'].mean(),inplace=True)

In [26]: df.isna().sum()

Out[26]: male                0
age                0
education          0
currentSmoker      0
cigsPerDay         0
BPMeds             0
prevalentStroke    0
prevalentHyp       0
diabetes           0
totChol            0
sysBP             0
diaBP             0
BMI               0
```

```
heartRate      0
glucose        0
TenYearCHD     0
dtype: int64
```

Logistic Regression Model

```
In [27]: #Splitting the dependent and independent variables.
x = df.drop("TenYearCHD",axis=1)
y = df['TenYearCHD']
```

```
In [28]: x #checking the features
```

```
Out[28]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	d
0	1	39	4.0	0	0.0	0.00000	0	0	
1	0	46	2.0	0	0.0	0.00000	0	0	
2	1	48	1.0	1	20.0	0.00000	0	0	
3	0	61	3.0	1	30.0	0.00000	0	1	
4	0	46	3.0	1	23.0	0.00000	0	0	
...
4233	1	50	1.0	1	1.0	0.00000	0	1	
4234	1	51	3.0	1	43.0	0.00000	0	0	
4235	0	48	2.0	1	20.0	0.02963	0	0	
4236	0	44	1.0	1	15.0	0.00000	0	0	
4237	0	52	2.0	0	0.0	0.00000	0	0	

4238 rows × 15 columns

Train Test Split

```
In [29]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_stat
```

```
In [30]: y_train
```

```
Out[30]:
```

3252	0
3946	0
1261	0
2536	0
4089	0
..	
3444	0
466	0
3092	0
3772	0
860	0

Name: TenYearCHD, Length: 3390, dtype: int64

KNN Classifier

```
In [31]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=5, p=2, metric='minkowski')
knn.fit(x_train, y_train)
acc = knn.score(x_test, y_test)*100
print(acc)
```

83.13679245283019

In []: